

STEP_5_LLAMA_DATASET_API.md

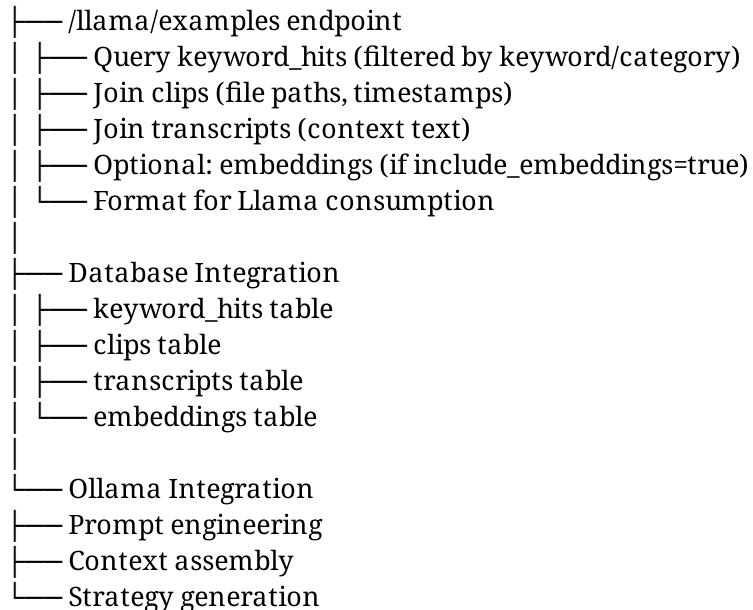
Step 5: Llama Dataset API Extension - Complete Implementation Guide

Overview

Step 5 enhances the /llama/examples endpoint with real database integration, embeddings support, and comprehensive Ollama integration for RAG-style workflows.

Architecture

API Service (Port 8003)



Implementation Tasks

Task 1: Database Integration for /llama/examples

Location: services/api-service/app/main.py

Replace the placeholder /llama/examples endpoint with full database implementation:

```
from sqlalchemy.orm import Session
from app.database import get_db
from app.models import KeywordHit, Clip, Transcript, Embedding, MediaItem
from fastapi import Depends
```

```
@app.get("/llama/examples", tags=["Llama Dataset"])
async def get_llama_examples(
    keyword: Optional[str] = Query(None),
    top_k: int = Query(5, ge=1, le=50),
    include_embeddings: bool = Query(False),
    category: Optional[str] = Query(None),
    db: Session = Depends(get_db)
):
    """

```

Get structured training examples for Llama fine-tuning or prompting

Query Parameters:

- keyword: Filter by keyword (e.g., "RSI", "breakout")
- top_k: Number of examples to return
- include_embeddings: Whether to include embedding vectors
- category: Filter by keyword category (technical_indicator, price_action, etc.)

Returns:

- List of examples with transcripts, clips, detected concepts, and optional embeddings

try:

```
# Build query
query = db.query(KeywordHit)
```

```
if keyword:
```

```
    query = query.filter(KeywordHit.keyword.ilike(f"%{keyword}%"))
```

```
if category:
```

```
    query = query.filter(KeywordHit.category == category)
```

```
# Order by confidence and limit
```

```
keyword_hits = query.order_by(KeywordHit.confidence.desc()).limit(top_k).all()
```

```
examples = []
```

```
for hit in keyword_hits:
```

```
    # Get associated clip
```

```
    clip = db.query(Clip).filter(Clip.keyword_hit_id == hit.id).first()
```

```
    # Get transcript context (segments around the hit)
```

```

transcript_segments = db.query(Transcript).filter(
    Transcript.media_item_id == hit.media_item_id,
    Transcript.start_time >= hit.start_time - 10, # 10 seconds before
    Transcript.end_time <= hit.end_time + 10      # 10 seconds after
).order_by(Transcript.segment_index).all()

transcript_text = " ".join([seg.text for seg in transcript_segments])

# Build example
example = {
    "clip_id": f"media_{hit.media_item_id}_{hit.keyword}_{hit.id}",
    "transcript": transcript_text,
    "keyword": hit.keyword,
    "category": hit.category if hasattr(hit, 'category') else None,
    "timestamp": hit.start_time,
    "clip_url": f"/clip/{clip.id}/download" if clip else None,
    "detected_concepts": [], # TODO: Add from ML service
    "context_text": hit.context_text,
    "confidence": hit.confidence
}

# Add frame path if available
if clip:
    example["frame_path"] = f"/data/processed/clips/media_{hit.media_item_id}_{hit.keyword}_{hit.id}.mp4"
    example["start_time"] = clip.start_time
    example["end_time"] = clip.end_time

# Add embeddings if requested
if include_embeddings:
    embedding = db.query(Embedding).filter(
        Embedding.media_item_id == hit.media_item_id,
        Embedding.reference_id == hit.id,
        Embedding.embedding_type == 'keyword'
    ).first()

    if embedding:
        example["embeddings"] = embedding.embedding_vector
        example["embedding_model"] = embedding.embedding_model

```

```

        else:
            example["embeddings"] = None

        examples.append(example)

    return {
        "examples": examples,
        "total": len(examples),
        "keyword_filter": keyword,
        "category_filter": category
    }

except Exception as e:
    logger.error(f"Error getting Llama examples: {str(e)}")
    raise HTTPException(status_code=500, detail=str(e))

```

Task 2: Create Llama Integration Module

Create new file: services/api-service/app/llama_client.py

services/api-service/app/llama_client.py

Client for interacting with Ollama for strategy generation

```

import requests
import logging
from typing import List, Dict, Optional

logger = logging.getLogger(name)

class LlamaClient:
    """Client for Ollama API integration"""

    def __init__(self, base_url: str = "http://localhost:11434"):
        self.base_url = base_url
        self.api_url = f"{base_url}/api/generate"

    def generate_strategy_from_examples(
        self,

```

```
examples: List[Dict],
model: str = "llama2",
temperature: float = 0.7
) -> Dict:
"""
Generate trading strategy from Llama examples

Args:
examples: List of keyword examples with transcripts
model: Ollama model name
temperature: Generation temperature (0-1)

Returns:
Dict with generated strategy and metadata
"""

try:
# Build context from examples
context = self._build_context(examples)

# Build prompt
prompt = self._build_strategy_prompt(context, examples)

# Call Ollama
response = requests.post(
    self.api_url,
    json={
        "model": model,
        "prompt": prompt,
        "stream": False,
        "options": {
            "temperature": temperature,
            "num_predict": 500
        }
    },
    timeout=60
)

if response.status_code != 200:
```

```

        logger.error(f"Ollama error: {response.text}")
        return {"status": "failed", "error": response.text}

    result = response.json()

    return {
        "status": "success",
        "strategy": result.get("response", ""),
        "model": model,
        "context_examples": len(examples),
        "prompt_length": len(prompt)
    }

except Exception as e:
    logger.error(f"Error generating strategy: {str(e)}")
    return {"status": "failed", "error": str(e)}

def _build_context(self, examples: List[Dict]) -> str:
    """Build context string from examples"""
    context_parts = []

    for i, ex in enumerate(examples, 1):
        context_parts.append(
            f"{i}. {ex['keyword']} at {ex['timestamp']}s:\n"
            f"  {ex['transcript'][:200]}...\n"
        )

    return "\n".join(context_parts)

def _build_strategy_prompt(self, context: str, examples: List[Dict]) -> str:
    """Build strategy generation prompt"""
    keywords = list(set([ex['keyword'] for ex in examples]))

    prompt = f"""You are a trading strategy expert. Based on the following video

```

Video Excerpts:
{context}

Key Concepts: ','.join(keywords)}

Please provide:

1. Strategy Overview (2-3 sentences)
2. Entry Conditions (specific indicators and levels)
3. Exit Conditions (take profit and stop loss)
4. Risk Management (position sizing, max drawdown)
5. Timeframe (best timeframe for this strategy)

Be specific with indicator values and thresholds where mentioned in the excerpts.""""

```
return prompt
```

```
def summarize_concept(  
    self,  
    keyword: str,  
    examples: List[Dict],  
    model: str = "llama2"  
) -> str:  
    """
```

Summarize a trading concept from examples

Args:

```
    keyword: The trading keyword/concept  
    examples: Examples containing the keyword  
    model: Ollama model name
```

Returns:

```
    Summary text
```

```
    """
```

```
try:
```

```
    context = self._build_context(examples)
```

```
    prompt = f"""Based on these trading video excerpts about {keyword}, prov
```

Excerpts:

```
{context}
```

Explanation:""""

```
response = requests.post(  
    self.api_url,
```

```

        json={
            "model": model,
            "prompt": prompt,
            "stream": False,
            "options": {"temperature": 0.5, "num_predict": 150}
        },
        timeout=30
    )

    if response.status_code == 200:
        return response.json().get("response", "")
    else:
        return f"Error: {response.text}"

except Exception as e:
    logger.error(f"Error summarizing concept: {str(e)}")
    return f"Error: {str(e)}"

```

Task 3: Add Strategy Generation Endpoint

Add to: services/api-service/app/main.py

```
from app.llama_client import LlamaClient
```

Initialize Llama client on startup

```

llama_client: Optional[LlamaClient] = None

@app.on_event("startup")
async def startup_event():
    """Initialize clients on startup"""
    global llama_client
    try:
        ollama_url = os.getenv("OLLAMA_URL", "http://localhost:11434")
        llama_client = LlamaClient(base_url=ollama_url)
        logger.info("Llama client initialized")
    except Exception as e:
        logger.error(f"Error initializing Llama client: {str(e)})"

@app.post("/llama/generate-strategy", tags=["Llama Dataset"])
async def generate_strategy_from_keyword(
    keyword: str,
    top_k: int = Query(5, ge=1, le=20),
    model: str = Query("llama2"),
    temperature: float = Query(0.7, ge=0.0, le=1.0),

```

```
db: Session = Depends(get_db)
):
"""

```

```
Generate a trading strategy using Llama based on keyword examples
```

Query Parameters:

- keyword: Trading keyword to base strategy on (e.g., "RSI", "MACD")
- top_k: Number of examples to use as context
- model: Ollama model name (default: llama2)
- temperature: Generation temperature (0-1)

Returns:

- Generated strategy with entry/exit conditions and risk management

```
"""

```

```
try:
```

```
    if llama_client is None:
```

```
        raise HTTPException(status_code=503, detail="Llama client not initialized")
```

```
# Get examples
```

```
examples_response = await get_llama_examples(
    keyword=keyword,
    top_k=top_k,
    include_embeddings=False,
    category=None,
    db=db
)
```

```
examples = examples_response["examples"]
```

```
if not examples:
```

```
    raise HTTPException(status_code=404, detail=f"No examples found for key
```

```
# Generate strategy
```

```
result = llama_client.generate_strategy_from_examples(
    examples=examples,
    model=model,
    temperature=temperature
)
```

```

if result["status"] == "failed":
    raise HTTPException(status_code=500, detail=result["error"])

return {
    "keyword": keyword,
    "strategy": result["strategy"],
    "model": result["model"],
    "examples_used": result["context_examples"],
    "timestamp": datetime.utcnow().isoformat()
}

except HTTPException:
    raise
except Exception as e:
    logger.error(f"Error generating strategy: {str(e)}")
    raise HTTPException(status_code=500, detail=str(e))

```

```

@app.get("/llama/summarize/{keyword}", tags=["Llama Dataset"])
async def summarize_keyword_concept(
    keyword: str,
    top_k: int = Query(3, ge=1, le=10),
    db: Session = Depends(get_db)
):
    """
    Get a Llama-generated summary of a trading concept

```

Path Parameters:

- keyword: Trading keyword to summarize

Query Parameters:

- top_k: Number of examples to base summary on

Returns:

- Concise explanation of the concept

"""

try:

- if llama_client is None:

- raise HTTPException(status_code=503, detail="Llama client not initialized")

```
# Get examples
examples_response = await get_llama_examples(
    keyword=keyword,
    top_k=top_k,
    include_embeddings=False,
    category=None,
    db=db
)

examples = examples_response["examples"]

if not examples:
    raise HTTPException(status_code=404, detail=f"No examples found for key

# Generate summary
summary = llama_client.summarize_concept(
    keyword=keyword,
    examples=examples
)

return {
    "keyword": keyword,
    "summary": summary,
    "examples_used": len(examples),
    "timestamp": datetime.utcnow().isoformat()
}

except HTTPException:
    raise
except Exception as e:
    logger.error(f"Error summarizing concept: {str(e)}")
    raise HTTPException(status_code=500, detail=str(e))
```

Task 4: Create Python Integration Examples

Create new file: docs/LLAMA_INTEGRATION_EXAMPLES.md

Llama Integration Examples

Python Script for Strategy Generation

```
import requests
import json
```

Configuration

```
API_BASE = "http://localhost:8003"
KEYWORD = "RSI"
```

1. Get examples for the keyword

```
examples_response = requests.get(
f"{API_BASE}/llama/examples",
params={
"keyword": KEYWORD,
"top_k": 5,
"include_embeddings": False
})
examples = examples_response.json()
print(f"Found {examples['total']} examples for '{KEYWORD}'")
```

2. Generate strategy using Llama

```
strategy_response = requests.post(
f"{API_BASE}/llama/generate-strategy",
params={
"keyword": KEYWORD,
"top_k": 5,
"model": "llama2",
"temperature": 0.7
})
strategy = strategy_response.json()
print("\n" + "*60)
print(f"TRADING STRATEGY FOR {KEYWORD}")
print("*60)
print(strategy['strategy'])
print("*60)
```

3. Get concept summary

```
summary_response = requests.get(  
f"{API_BASE}/llama/summarize/{KEYWORD}",  
params={"top_k": 3}  
)  
  
summary = summary_response.json()  
print(f"\nConcept Summary: {summary['summary']}")
```

Direct Ollama Integration

```
import requests
```

Get examples from API

```
examples_response = requests.get(  
"http://localhost:8003/llama/examples",  
params={"keyword": "MACD", "top_k": 5}  
)  
  
examples = examples_response.json()["examples"]
```

Build custom prompt

```
context = "\n".join([  
f"- {ex['transcript'][:150]}... (at {ex['timestamp']}s)"  
for ex in examples  
])  
  
prompt = f"""Based on these trading video excerpts about MACD:  
{context}"""
```

Generate a step-by-step trading strategy using MACD. Include:

1. Entry signals
2. Exit signals
3. Stop loss placement
4. Position sizing""""

Send to Ollama directly

```
ollama_response = requests.post(  
"http://localhost:11434/api/generate",  
json={  
"model": "llama2",  
"prompt": prompt,  
"stream": False
```

```
}

strategy = ollama_response.json()["response"]
print(strategy)
```

Batch Processing Multiple Keywords

```
import requests
from concurrent.futures import ThreadPoolExecutor

API_BASE = "http://localhost:8003"
KEYWORDS = ["RSI", "MACD", "Bollinger Bands", "Support", "Resistance"]

def generate_strategy(keyword):
    """Generate strategy for a keyword"""
    response = requests.post(
        f"{API_BASE}/llama/generate-strategy",
        params={"keyword": keyword, "top_k": 5}
    )
    return keyword, response.json()
```

Process in parallel

```
with ThreadPoolExecutor(max_workers=3) as executor:
    results = executor.map(generate_strategy, KEYWORDS)

for keyword, result in results:
    print(f"\n{'='*60}")
    print(f"STRATEGY: {keyword}")
    print('='*60)
    print(result['strategy'])
```

Task 5: Update API Documentation

Update: docs/API_DOCUMENTATION.md

Add new section after the existing Llama Dataset API section:

Llama Strategy Generation (NEW)

POST /llama/generate-strategy

Generate a complete trading strategy using Llama based on keyword examples

Query Parameters:

- keyword (string, required): Trading keyword (e.g., "RSI", "breakout")
- top_k (int, default: 5): Number of examples to use as context
- model (string, default: "llama2"): Ollama model name
- temperature (float, default: 0.7): Generation temperature (0-1)

Example:

```
curl -X POST "http://localhost:8003/llama/generate-strategy?keyword=RSI&top\_k=5&model=llama2&temperature=0.7"
```

Response:

```
{  
  "keyword": "RSI",  
  "strategy": "Strategy Overview:\nThis RSI-based momentum strategy...\n\nEntry Conditions:\n1. RSI crosses above 30 (oversold)\n2. Price confirms with bullish candle...",  
  "model": "llama2",  
  "examples_used": 5,  
  "timestamp": "2025-11-20T21:00:00Z"  
}
```

GET /llama/summarize/{keyword}

Get a concise Llama-generated summary of a trading concept

Path Parameters:

- keyword (string): Trading keyword to summarize

Query Parameters:

- top_k (int, default: 3): Number of examples to base summary on

Example:

```
curl "http://localhost:8003/llama/summarize/MACD?top\_k=3"
```

Response:

```
{  
  "keyword": "MACD",  
  "summary": "MACD (Moving Average Convergence Divergence) is a momentum indicator that shows the relationship between two moving averages. It generates buy signals when the MACD line crosses above the signal line and sell signals when it crosses below. Traders also watch for divergences between MACD and price action as potential reversal signals.",  
  "examples_used": 3,  
  "timestamp": "2025-11-20T21:00:00Z"  
}
```

Testing Step 5

End-to-End Test Script

Create: tests/test_llama_integration.py

```
import requests  
import time  
  
API_BASE = "http://localhost:8003"  
  
def test_llama_pipeline():  
    """Test complete Llama integration pipeline"""
```

```

print("Step 1: Testing /llama/examples endpoint...")
examples_response = requests.get(
    f"{API_BASE}/llama/examples",
    params={"keyword": "RSI", "top_k": 3}
)
assert examples_response.status_code == 200
examples = examples_response.json()
print(f"✓ Found {examples['total']} examples")

print("\nStep 2: Testing /llama/generate-strategy endpoint...")
strategy_response = requests.post(
    f"{API_BASE}/llama/generate-strategy",
    params={"keyword": "RSI", "top_k": 3}
)
assert strategy_response.status_code == 200
strategy = strategy_response.json()
print(f"✓ Generated strategy ({len(strategy['strategy'])} chars)")

print("\nStep 3: Testing /llama/summarize endpoint...")
summary_response = requests.get(
    f"{API_BASE}/llama/summarize/RSI",
    params={"top_k": 2}
)
assert summary_response.status_code == 200
summary = summary_response.json()
print(f"✓ Generated summary: {summary['summary'][:100]}...")

print("\n✓ All Llama integration tests passed!")

```

```

if name == 'main':
    test_llama_pipeline()

```

Implementation Checklist

- [] Update /llama/examples endpoint with database integration
- [] Create llama_client.py module
- [] Add LlamaClient initialization in startup event
- [] Implement /llama/generate-strategy endpoint
- [] Implement /llama/summarize/{keyword} endpoint

- [] Create Python integration examples documentation
 - [] Update API documentation with new endpoints
 - [] Create end-to-end test script
 - [] Test with sample data
 - [] Verify Ollama connectivity
-

Configuration

Update .env file:

Ollama Configuration

```
OLLAMA_URL=http://ollama:11434  
OLLAMA_MODEL=llama2  
OLLAMA_TEMPERATURE=0.7
```

Next Steps

After Step 5 completion:

Step 6: Strategy & Backtesting Framework

- Feature engineering from embeddings
- ML/RL model training
- Enhanced backtesting service
- Strategy promotion logic

Step 7: Tests & Acceptance

- Unit tests for all endpoints
 - Integration tests
 - Performance benchmarks
 - CI/CD pipeline
-

Status

✓ Step 5 Ready for Implementation - Llama Dataset API Extension

- Database integration patterns defined
- LlamaClient module designed
- 3 new endpoints specified
- Integration examples created
- Testing approach documented

Implementation Time Estimate: 2-3 hours